

Lab 1

Frank Palma Gomez

11:59PM February 18, 2021

You should have RStudio installed to edit this file. You will write code in places marked “TO-DO” to complete the problems. Some of this will be a pure programming assignment. The tools for the solutions to these problems can be found in the class practice lectures. I want you to use the methods I taught you, not for you to google and come up with whatever works. You won’t learn that way.

To “hand in” the homework, you should compile or publish this file into a PDF that includes output of your code. Once it’s done, push by the deadline to your repository in a directory called “labs”.

- Print out the numerical constant pi with ten digits after the decimal point using the internal constant pi.

```
options(digits=11) # exclusive this is 10
pi
```

```
## [1] 3.1415926536
```

- Sum up the first 103 terms of the series $1 + 1/2 + 1/4 + 1/8 + \dots$

```
sum(1 / (2^(0:102)))
```

```
## [1] 2
```

- Find the product of the first 37 terms in the sequence $1/3, 1/6, 1/9 \dots$

```
prod(1 / (seq(from=3, by=3, length.out=37)))
```

```
## [1] 1.613528728e-61
```

- Find the product of the first 387 terms of $1 * 1/2 * 1/4 * 1/8 * \dots$

```
prod(1 / 2^(0:386)) # underflow
```

```
## [1] 0
```

Is this answer *exactly* correct?

Not exactly correct because we reached the smallest possible number. We experienced numerical underflow.

- Figure out a means to express the answer more exactly. Not compute exactly, but express more exactly.

```
-log(2)*sum(0:386)
```

```
## [1] -51771.856063
```

- Create the sequence $x = [\text{Inf}, 20, 18, \dots, -20]$.

```
x <- c(Inf, seq(from=20, to=-20, by=-2))
x
```

```
## [1] Inf 20 18 16 14 12 10 8 6 4 2 0 -2 -4 -6 -8 -10 -12 -14
## [20] -16 -18 -20
```

Create the sequence `x = [log3(Inf), log3(100), log3(98), ... log3(-20)]`.

```
x <- c(Inf, seq(from=100, to=-20, by=-2))
x = log(x, base=3)
```

```
## Warning: NaNs produced
```

Comment on the appropriateness of the non-numeric values.

There is a -Inf because `log(0)` is -Inf. The proceeding are NaN because `log` is not defined for negative answers

- Create a vector of booleans where the entry is true if `x[i]` is positive and finite.

```
y = !is.nan(x) & is.finite(x) & x > 0
```

- Locate the indices of the non-real numbers in this vector. Hint: use the `which` function. Don't hesitate to use the documentation via `?which`.

```
which(y == FALSE)
```

```
## [1] 1 52 53 54 55 56 57 58 59 60 61 62
```

- Locate the indices of the infinite quantities in this vector.

```
which(is.infinite(x))
```

```
## [1] 1 52
```

- Locate the indices of the min and max in this vector. Hint: use the `which.min` and `which.max` functions.

```
which.min(x)
```

```
## [1] 52
```

```
which.max(x)
```

```
## [1] 1
```

- Count the number of unique values in `x`.

```
length(unique(x))
```

```
## [1] 53
```

- Cast `x` to a factor. Do the number of levels make sense?

```
as.factor(x)
```

```
## [1] Inf
## [5] 4.13548512895119 4.11590933734319 4.09590327428938 4.07544759935851
## [9] 4.05452163806914 4.03310325630434 4.01116871959141 3.98869253500376
## [13] 3.96564727304425 3.94200336638929 3.91772888178973 3.89278926071437
## [17] 3.86714702345081 3.84076143030548 3.81358809221559 3.78557852142874
## [21] 3.75667961082847 3.72683302786084 3.69597450568212 3.66403300987579
## [25] 3.63092975357146 3.59657702661571 3.56087679500731 3.52371901428583
## [29] 3.48497958377173 3.44451784578705 3.40217350273288 3.3577627814323
## [33] 3.31107361281783 3.26185950714291 3.20983167673402 3.15464876785729
## [37] 3.09590327428938 3.03310325630434 2.96564727304425 2.89278926071437
## [41] 2.8135880922156 2.72683302786084 2.63092975357146 2.52371901428583
## [45] 2.40217350273288 2.26185950714291 2.09590327428938 1.89278926071437
## [49] 1.63092975357146 1.26185950714291 0.630929753571457 -Inf
## [53] NaN NaN NaN NaN
## [57] NaN NaN NaN NaN
## [61] NaN NaN
```

```
## 53 Levels: -Inf 0.630929753571457 1.26185950714291 ... NaN
```

- Cast `x` to integers. What do we learn about R's infinity representation in the integer data type?

NaN got converted to NA

```
as.integer(x)
```

```
## Warning: NAs introduced by coercion to integer range
```

```
## [1] NA 4 4 4 4 4 4 4 4 4 4 3 3 3 3 3 3 3 3 3 3 3 3
## [26] 3 3 3 3 3 3 3 3 3 3 3 3 3 2 2 2 2 2 2 2 2 1 1 1
## [51] 0 NA NA NA NA NA NA NA NA NA NA NA NA NA
```

- Use `x` to create a new vector `y` containing only the real numbers in `x`.

```
y = x[!is.nan(x) & is.finite(x) & x > 0]
y
```

```
## [1] 4.19180654858 4.17341725189 4.15464876786 4.13548512895 4.11590933734
## [6] 4.09590327429 4.07544759936 4.05452163807 4.03310325630 4.01116871959
## [11] 3.98869253500 3.96564727304 3.94200336639 3.91772888179 3.89278926071
## [16] 3.86714702345 3.84076143031 3.81358809222 3.78557852143 3.75667961083
## [21] 3.72683302786 3.69597450568 3.66403300988 3.63092975357 3.59657702662
## [26] 3.56087679501 3.52371901429 3.48497958377 3.44451784579 3.40217350273
## [31] 3.35776278143 3.31107361282 3.26185950714 3.20983167673 3.15464876786
## [36] 3.09590327429 3.03310325630 2.96564727304 2.89278926071 2.81358809222
## [41] 2.72683302786 2.63092975357 2.52371901429 2.40217350273 2.26185950714
## [46] 2.09590327429 1.89278926071 1.63092975357 1.26185950714 0.63092975357
```

- Use the left rectangle method to numerically integrate x^2 from 0 to 1 with rectangle width size $1e-6$.

```
sum(seq(from=0, to=1 - 1e-6, by=1e-6)^2) * 1e-6
```

```
## [1] 0.33333283333
```

- Calculate the average of 100 realizations of standard Bernoullis in one line using the `sample` function.

```
sample(c(0, 1), size=100, replace=TRUE)
```

```
## [1] 1 1 1 1 0 0 1 1 0 0 1 1 0 0 1 0 1 1 1 0 1 1 0 1 0 0 0 1 1 0 0 1 1 0 0 0
## [38] 1 1 1 0 1 0 0 1 0 0 1 1 0 1 1 0 1 0 1 0 0 0 1 1 0 0 1 0 0 1 0 1 0 1 1 1 1
## [75] 0 1 0 1 0 1 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0
```

- Calculate the average of 500 realizations of Bernoullis with $p = 0.9$ in one line using the `sample` and `mean` functions.

```
sample(c(0, 1), size=500, replace=TRUE, prob=c(0.1, 0.9))
```

```
## [1] 1 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1 1 1 0 1 1 1 1 1 0 1 0 0 1 0 1 1 1 1 1 1
## [38] 1 1 1 1 1 1 1 1 0 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1
## [75] 0 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1
## [112] 1 0 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0
## [149] 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [186] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0
## [223] 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [260] 1 1 1 1 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1
## [297] 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [334] 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 0 0 1 1 1 1 1 1 1 1 1
## [371] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1
## [408] 1 1 1 0 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 1 1 1 0 1 1 1 1 1 1 1 1
```

```
## [445] 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1
## [482] 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

- Calculate the average of 1000 realizations of Bernoullis with $p = 0.9$ in one line using `rbinom`.

```
mean(rbinom(n=1000, size=1, prob=0.9))
```

```
## [1] 0.909
```

- In class we considered a variable `x_3` which measured “criminality”. We imagined $L = 4$ levels “none”, “infraction”, “misdemeanor” and “felony”. Create a variable `x_3` here with 100 random elements (equally probable). Create it as a nominal (i.e. unordered) factor.

```
x_3 = as.factor(sample(c("none", "infraction", "misdemeanor", "felony"), size=100, replace=TRUE))
```

- Use `x_3` to create `x_3_bin`, a binary feature where 0 is no crime and 1 is any crime.

```
x_3_bin = x_3 != "none"
x_3_bin
```

```
## [1] TRUE TRUE FALSE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [13] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE
## [25] TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE
## [37] TRUE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [49] TRUE TRUE FALSE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE
## [61] TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE TRUE
## [73] TRUE TRUE FALSE TRUE TRUE FALSE FALSE TRUE TRUE FALSE TRUE FALSE
## [85] FALSE FALSE TRUE TRUE FALSE TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE
## [97] FALSE FALSE TRUE TRUE
```

- Use `x_3` to create `x_3_ord`, an ordered factor variable. Ensure the proper ordinal ordering.

```
x_3_ord = factor(x_3, levels=c("none", "infraction", "misdemeanor", "felony"), order=TRUE)
x_3_ord
```

```
## [1] felony felony none none felony misdemeanor
## [7] felony felony misdemeanor misdemeanor felony misdemeanor
## [13] felony misdemeanor felony felony felony infraction
## [19] infraction misdemeanor infraction misdemeanor none felony
## [25] misdemeanor misdemeanor infraction felony misdemeanor infraction
## [31] misdemeanor infraction felony infraction none felony
## [37] infraction infraction none infraction none felony
## [43] infraction felony infraction felony misdemeanor infraction
## [49] felony felony none infraction infraction infraction infraction
## [55] none felony infraction infraction infraction misdemeanor
## [61] misdemeanor misdemeanor misdemeanor felony felony none
## [67] infraction felony infraction felony felony felony
## [73] infraction misdemeanor none infraction infraction none
## [79] none misdemeanor infraction none misdemeanor none
## [85] none none misdemeanor misdemeanor none misdemeanor
## [91] misdemeanor infraction infraction infraction none infraction
## [97] none none infraction felony
## Levels: none < infraction < misdemeanor < felony
```

- Convert this variable into three binary variables without any information loss and put them into a data matrix.

```
x_3_infraction = as.integer(x_3 == "infraction")
x_3_misdemeanor = as.integer(x_3 == "misdemeanor")
```

```
x_3_felony = as.integer(x_3 == "felony")

X = cbind(x_3_infraction, x_3_misdemeanor, x_3_felony) # concatenate column wise

cols = levels(x_3)[1:3] # get the first three columns
X = matrix(X, nrow=length(x_3), ncol=length(cols)) # convert into matrix
colnames(X) = cols # set column names

head(X) # print
```

```
##      felony infraction misdemeanor
## [1,]      0          0            1
## [2,]      0          0            1
## [3,]      0          0            0
## [4,]      0          0            0
## [5,]      0          0            1
## [6,]      0          0            0
```

- What should the sum of each row be (in English)?

The sum of each row should sum crimes an observation has. It should result in a vector.

Verify that.

```
row_counts = rowSums(X, dims=1)
row_counts

##      [1] 1 1 0 0 1 0 1 1 0 0 1 0 1 0 1 1 1 1 0 1 0 0 1 0 0 1 1 0 1 0 1 1 1 0 1 1
##      [38] 1 0 1 0 1 1 1 1 1 0 1 1 1 0 1 1 1 0 1 1 1 1 0 0 0 0 1 1 0 1 1 1 1 1 1 1 0
##      [75] 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 1 0 1 0 0 1 1
```

- How should the column sum look (in English)?

It should be a vector with the size of the number of columns

Verify that.

```
col_counts = colSums(X, dims=1)
col_counts

##      felony  infraction misdemeanor
##          30           0            27
```

- Generate a matrix with 100 rows where the first column is realization from a normal with mean 17 and variance 38, the second column is uniform between -10 and 10, the third column is poisson with mean 6, the fourth column in exponential with lambda of 9, the fifth column is binomial with $n = 20$ and $p = 0.12$ and the sixth column is a binary variable with exactly 24% 1's dispersed randomly. Name the rows the entries of the `fake_first_names` vector.

```
fake_first_names = c(
  "Sophia", "Emma", "Olivia", "Ava", "Mia", "Isabella", "Riley",
  "Aria", "Zoe", "Charlotte", "Lily", "Layla", "Amelia", "Emily",
  "Madelyn", "Aubrey", "Adalyn", "Madison", "Chloe", "Harper",
  "Abigail", "Aaliyah", "Avery", "Evelyn", "Kaylee", "Ella", "Ellie",
  "Scarlett", "Arianna", "Hailey", "Nora", "Addison", "Brooklyn",
  "Hannah", "Mila", "Leah", "Elizabeth", "Sarah", "Eliana", "Mackenzie",
  "Peyton", "Maria", "Grace", "Adeline", "Elena", "Anna", "Victoria",
  "Camilla", "Lillian", "Natalie", "Jackson", "Aiden", "Lucas",
  "Liam", "Noah", "Ethan", "Mason", "Caden", "Oliver", "Elijah",
```

```

"Grayson", "Jacob", "Michael", "Benjamin", "Carter", "James",
"Jayden", "Logan", "Alexander", "Caleb", "Ryan", "Luke", "Daniel",
"Jack", "William", "Owen", "Gabriel", "Matthew", "Connor", "Jayce",
"Isaac", "Sebastian", "Henry", "Muhammad", "Cameron", "Wyatt",
"Dylan", "Nathan", "Nicholas", "Julian", "Eli", "Levi", "Isaiah",
"Landon", "David", "Christian", "Andrew", "Brayden", "John",
"Lincoln"
)
norm = rnorm(n=100, mean=17, sd=sqrt(38))
unif = runif(n=100, min=-10, max=10)
pois = rpois(n=100, lambda=6)
exp = rexp(n=100, rate=9)
binom = rbinom(n=20, size=1, p=0.12)
rand = rbinom(n=100, size=1, p=0.24)
cols = c('norm', 'unif', 'pois', 'exp', 'binom', 'binary')

X = cbind(norm, unif, pois, exp, binom, rand)
X = matrix(X, nrow=100, ncol=6)
rownames(X) = fake_first_names
colnames(X) = cols
head(X)

```

```

##           norm           unif pois           exp binom binary
## Sophia  20.580742010  4.42339935806   8 0.0224850507867    0    0
## Emma    15.673824956 -0.32417375594   8 0.0618617391317    0    0
## Olivia  20.272108289 -6.54996076133   4 0.0369593734439    0    1
## Ava     14.714160775  2.68836348318   3 0.0919946807700    0    0
## Mia     21.003815417 -8.07490918785   8 0.1427429807702    0    0
## Isabella 22.800037966  8.09769294690   9 0.0040534077659    0    0

```

- Create a data frame of the same data as above except make the binary variable a factor “DOMESTIC” vs “FOREIGN” for 0 and 1 respectively. Use RStudio’s View function to ensure this worked as desired.

```

X = data.frame(
  norm = X[, "norm"],
  unif = X[, "unif"],
  pois = X[, "pois"],
  exp = X[, "exp"],
  binom = X[, "binom"],
  binary = X[, "binary"]
)
X$binary = factor(X$binary, labels=c("DOMESTIC", "FOREIGN"))

```

- Print out a table of the binary variable. Then print out the proportions of “DOMESTIC” vs “FOREIGN”.

```
table(X$binary)
```

```

##
## DOMESTIC FOREIGN
##          76      24

```

Print out a summary of the whole dataframe.

```
summary(X)
```

```

##           norm           unif           pois
## Min.      : 6.5042103   Min.      : -9.79503158   Min.      : 1.00

```

```
## 1st Qu.:12.9480580 1st Qu.: -5.23529199 1st Qu.: 4.00
## Median :16.6386898 Median : -0.24723901 Median : 6.00
## Mean :17.2579218 Mean : -0.54746954 Mean : 5.81
## 3rd Qu.:21.4884514 3rd Qu.: 3.96955856 3rd Qu.: 7.00
## Max. :33.6345558 Max. : 9.56398258 Max. :13.00
## exp binom binary
## Min. :0.00004730085 Min. :0 DOMESTIC:76
## 1st Qu.:0.03038306949 1st Qu.:0 FOREIGN :24
## Median :0.06856002064 Median :0
## Mean :0.11191047990 Mean :0
## 3rd Qu.:0.14840337885 3rd Qu.:0
## Max. :0.59275266711 Max. :0
```

- Let $n = 50$. Create a $n \times n$ matrix R of exactly 50% entries 0's, 25% 1's 25% 2's. These values should be in random locations.

```
n = 50
R = matrix(sample(0:2, size=n * n, prob=c(0.5, 0.25, 0.25), replace=TRUE), nrow=n, ncol=n)
head(R)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
## [1,]  1    1    1    0    2    0    0    0    1    1    0    0    0    0
## [2,]  2    0    1    1    1    0    2    0    0    1    1    0    0    0
## [3,]  1    1    0    0    2    0    0    0    1    2    0    1    1    0
## [4,]  0    1    1    0    2    1    1    0    1    1    0    2    0    2
## [5,]  0    0    2    1    0    0    0    0    2    0    2    0    1    1
## [6,]  0    1    2    0    0    1    0    0    0    2    0    0    0    1
##      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
## [1,]  1    0    0    0    0    0    0    0    1    0    1    2    0
## [2,]  1    2    0    0    2    1    0    2    0    2    0    0    0
## [3,]  1    0    2    0    0    1    0    1    2    0    0    0    0
## [4,]  0    1    1    2    1    0    0    2    0    0    0    0    0
## [5,]  2    0    1    0    0    0    2    1    1    0    2    2    2
## [6,]  1    0    0    0    2    1    2    0    0    0    0    1    2
##      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
## [1,]  0    0    0    0    0    1    2    2    2    0    0    2
## [2,]  0    1    2    2    0    1    1    0    1    0    0    2
## [3,]  2    0    0    0    1    0    2    0    1    0    1    2
## [4,]  0    0    2    0    1    0    0    2    1    0    0    0
## [5,]  1    0    0    0    1    2    0    0    1    0    0    0
## [6,]  2    2    2    0    0    1    1    0    2    2    2    0
##      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
## [1,]  1    1    0    2    0    0    0    2    0    1    0    1
## [2,]  1    0    2    1    1    0    0    1    1    0    2    0
## [3,]  1    1    1    1    0    0    2    2    0    1    0    0
## [4,]  0    0    2    2    0    1    2    0    0    1    2    1
## [5,]  0    0    1    0    0    0    0    0    2    0    0    0
## [6,]  1    0    1    2    0    0    0    0    0    0    1    1
```

- Randomly punch holes (i.e. NA) values in this matrix so that an each entry is missing with probability 30%.

```
R = matrix(lapply(R, function(x) x[sample(c(TRUE, NA), size=length(x), prob=c(0.70, 0.30), replace=TRUE)]), nrow=n, ncol=n)
head(R)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14]
## [1,] NA    1    1    0    2    0    0    0    1    1    NA    0    NA    0
```

```
## [2,] 2 0 1 1 NA NA 2 0 0 NA NA NA 0 0
## [3,] 1 NA NA 0 2 0 0 0 NA NA 0 1 1 NA
## [4,] NA 1 1 0 2 1 1 NA 1 NA 0 NA NA 2
## [5,] NA 0 2 1 NA 0 0 0 2 0 NA 0 NA 1
## [6,] 0 1 NA 0 0 NA 0 NA 0 NA 0 0 0 1
##      [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26]
## [1,] 1 0 0 NA NA NA 0 1 0 1 2 0
## [2,] NA 2 NA 0 2 1 0 2 0 2 NA 0
## [3,] NA NA 2 0 0 NA 0 1 NA 0 0 0
## [4,] 0 1 1 2 1 NA NA NA 0 NA 0 0
## [5,] NA NA 1 0 0 0 2 1 NA NA NA 2
## [6,] 1 0 0 0 2 1 2 0 NA 0 1 2
##      [,27] [,28] [,29] [,30] [,31] [,32] [,33] [,34] [,35] [,36] [,37] [,38]
## [1,] 0 NA 0 NA 0 1 2 2 NA 0 0 2
## [2,] 0 NA NA 2 0 1 1 0 1 0 0 2
## [3,] 2 0 0 0 1 0 2 0 1 0 NA NA
## [4,] NA 0 NA 0 NA 0 0 NA 1 0 NA NA
## [5,] 1 0 0 0 1 2 NA 0 1 0 0 0
## [6,] NA 2 NA 0 0 1 NA 0 2 2 2 0
##      [,39] [,40] [,41] [,42] [,43] [,44] [,45] [,46] [,47] [,48] [,49] [,50]
## [1,] 1 1 NA 2 0 0 0 NA 0 1 0 1
## [2,] 1 NA 2 1 NA 0 0 NA 1 0 NA NA
## [3,] 1 1 1 1 0 0 2 2 NA 1 NA NA
## [4,] NA 0 2 2 NA NA NA 0 0 NA 2 1
## [5,] 0 NA NA 0 0 0 0 NA 2 0 0 NA
## [6,] 1 0 1 2 0 NA 0 0 0 NA 1 1
```

- Sort the rows in matrix R by the largest row sum to lowest. Be careful about the NA's!
- We will now learn the `apply` function. This is a handy function that saves writing for loops which should be eschewed in R. Use the `apply` function to compute a vector whose entries are the standard deviation of each row. Use the `apply` function to compute a vector whose entries are the standard deviation of each column. Be careful about the NA's! This should be one line.

```
#sd_fun = function(x) sd(na.exclude(x))
#R_row_sd = apply(R, 2, sd_fun) # sd for rows
#R_col_sd = apply(R, 2, sd_fun)
```

- Use the `apply` function to compute a vector whose entries are the count of entries that are 1 or 2 in each column. This should be one line.
- Use the `split` function to create a list whose keys are the column number and values are the vector of the columns. Look at the last example in the documentation `?split`.

```
?split
```

- In one statement, use the `lapply` function to create a list whose keys are the column number and values are themselves a list with keys: “min” whose value is the minimum of the column, “max” whose value is the maximum of the column, “pct_missing” is the proportion of missingness in the column and “first_NA” whose value is the row number of the first time the NA appears.

```
#T0-D0
```

- Set a seed and then create a vector v consisting of a sample of 1,000 iid normal realizations with mean -10 and variance 100.

```
set.seed(2000)
v = rnorm(1000, mean=-10, sd=sqrt(100))
```



```
head(v)
```

```
## [1] -18.5384323568 -13.5288939552 -1.0985047234  7.4128671643 -0.8392165958  
## [6] -21.7660142823
```

- Repeat this exercise by resetting the seed to ensure you obtain the same results.

```
set.seed(2000)
```

```
v = rnorm(1000, mean=-10, sd=sqrt(100))
```

```
head(v)
```

```
## [1] -18.5384323568 -13.5288939552 -1.0985047234  7.4128671643 -0.8392165958  
## [6] -21.7660142823
```

- Find the average of `v` and the standard error of `v`.

```
mean(v)
```

```
## [1] -9.4585200415
```

```
sd(v) / sqrt(length(v))
```

```
## [1] 0.31881861326
```

- Find the 5%ile of `v` and use the `qnorm` function to compute what it theoretically should be. Is the estimate about what is expected by theory?

```
quantile(v, 0.05)
```

```
##           5%
```

```
## -26.780618034
```

```
qnorm(0.05, mean=-10, sd=sqrt(100))
```

```
## [1] -26.44853627
```

- What is the percentile of `v` that corresponds to the value 0? What should it be theoretically? Is the estimate about what is expected by theory?

```
inverse_quant_obj = ecdf(v)
```

```
inverse_quant_obj(0)
```

```
## [1] 0.821
```